

STAT 534 Project

Spring 2019

Ashwin T. A. N


The aim of this project was to train a Hidden Markov Model (HMM) on the given training data and use it to analyze a test dataset. The training data set \mathcal{D} consisted of $K = 1000$ sequences of $T = 40$ observed states, with values in $\{0, 1, 2, 3\}$. The number of hidden states N was chosen by cross-validation and a best-fit Hidden Markov Model was computed using the Baum-Welch algorithm. Using this model, the most probable sequences of hidden states was computed for the given test data set \mathcal{E} of $K' = 50$ sequences by the Viterbi algorithm. Moreover, I also computed the probability distribution of the missing $(T + 1)^{st}$ observed states for each of the sequences in the test dataset. All analyses were implemented in Python version 2.7.

1 Training

The first step in training a HMM is to choose the number of hidden states N . This was done by cross-validation. The original training dataset was split into two parts - \mathcal{D}_1 consisting of the first $K_0 = 800$ sequences was used to train HMMs for different values of N by the Baum-Welch algorithm. These models were then tested on the set \mathcal{D}_2 consisting of the remaining $K - K_0 = 200$ sequences. Both the Akaike Information Criterion (AIC) and the Bayesian Information Criterion (BIC) indicated $N = 9$ to be the optimum choice.

Throughout the rest of this report, I will use A, B, π to represent the transitional probability matrix, the emission probability matrix and the initial distribution of hidden states for a HMM respectively. I started by writing modules that would, given a model (A, B, π) and a sequence of observed states O , compute the corresponding $\alpha_t(i), \beta_t(i)$ by the Forward-Backward algorithm. Using this module, it is easy to compute the log-likelihood of any given observed dataset.


Next, I implemented the Baum-Welch algorithm as follows: The model parameters A, B, π were initialized by choosing their entries independently from a Uniform(0, 1) distribution and then normalizing by rows. Let l_t denote the log-likelihood of \mathcal{D}_1 given the model output by the t -th iteration of the Baum-Welch algorithm. The iteration was made to repeat until $|l_{t+1}/l_t - 1| < 10^{-4}$ or the number of iterations reached 50. I then used this to train HMMs λ_N for $N = 5, 6, \dots, 10$. The log-likelihoods of the cross-validation data set \mathcal{D}_2 for these models is graphed below:



In addition, in order to account for the fact that the model increases in complexity as N increases, I also computed the Akaike Information Criterion (AIC) and the Bayesian Information Criterion (BIC) statistics. Recall that the AIC and BIC values are given by


$$\begin{aligned} AIC &= 2k - 2l, \\ BIC &= k \log(n) - 2l, \end{aligned}$$

where k is the number of parameters in the model, l is the log-likelihood and n is the sample size respectively. We note that after taking the row-sum constraints into account, the total number of real parameters in (A, B, π) is $k = N(N - 1) + N(4 - 1) + (N - 1) = N^2 + 3N - 1$. The resulting AIC, BIC values are graphed below:




All 3 of these plots suggest that choosing $N = 9$ is optimum in among $5, 6, \dots, 10$. Therefore,

with this choice of N , I ran the Baum-Welch algorithm on the entire training data set \mathcal{D} . The program ran for 50 iterations (the maximum allowed) with a total run time of 712 seconds. This implies an average run time of roughly 14.24 seconds per iteration. The following plot shows the growth of log-likelihood of \mathcal{D} with each iteration:



At the end of 50 iterations, the value of log-likelihood was -53192.25 . Since the total log-likelihood of the sample is equal to the sum of the log-likelihoods of each of the sequences in the sample, it is relevant to consider the ratio of log-likelihood to sample size, which gives the average log-likelihood of the sequences in the training data. Since our data had 1000 samples, this ratio turns out to be around -53.19 .


Finally, I relabeled the hidden states from $0, 1, \dots, 8$ such that the states with more uniform emission probabilities have a lower index. This was done by permuting the rows of A so that they are in ascending order of their l^1 norms, and then applying the same permutations to the columns of A , rows of B and π . The heatmaps of the resulting transmission matrix, emission matrix and initial distribution are graphed below:




2 Analysis of Test data

Let $\hat{\lambda} = (\hat{A}, \hat{B}, \hat{\pi})$ denote the Hidden Markov Model obtained in the previous section. The first task was to compute the log-likelihood of the test data \mathcal{E} under this model. The test data set given consisted of $K' = 50$ sequences, each one of length $T = 40$ as before. I computed the log-likelihood of this test data set to be -2719.53 . This yields a log-likelihood - to - sample size ratio of around -54.38 , which is quite close to -53.19 that we obtained for the training data set.

Next, I implemented the Viterbi algorithm with parameters $\hat{A}, \hat{B}, \hat{\pi}$ to compute the most likely sequence of hidden states corresponding to each sequence in the test data. The run-time of the algorithm was 0.074 seconds. The resulting output is illustrated in the plot below:



The states 6 and 7 make up a large majority of all the predicted states and most of the state transitions are between 6 and 7. 0 and 8 appear only once while 1 and 2 don't appear at all in the output. The output of the algorithm looked very similar for the given training data as well:



It would be interesting to see what happens for higher values of N .

Finally, I implemented an algorithm that computed the probability distribution of the missing $(T+1)^{th}$ observation in each of the sequences of the test data. Suppose $q_{1:T}$ is a sequence of hidden

states and $O_{1:T}$ the corresponding observed states generated by a HMM (A, B, π) . Then

$$\begin{aligned}
P[O_{T+1} = j | O_{1:T}] &= \sum_{i=1}^N \sum_{k=1}^N P[O_{T+1} = j, q_{T+1} = k, q_T = i | O_{1:T}] \\
&= \sum_{i=1}^N \sum_{k=1}^N \gamma_T(i) A_{ik} B_{kj} \\
&= (\gamma_T AB)(j).
\end{aligned}$$

I already had the Forward-Backward algorithm that computes $\alpha_t(i)$ and $\beta_t(i)$ for all t, i . From this, it is easy to compute $\gamma_T(\cdot)$ and further matrix multiplication by A and B gives us the required probability distribution. Iterating over all the sequences in the test data set now gives us the required output.